

Equal Weight per Instance Does Not Pay Alone: A Blob-Loss Auxiliary on Full-Resolution Cityscapes Buys Thin-Object IoU with Pedestrian Recall, and Only Pays as an Expert of a Mixture-of-Experts

Guillaume Cassez · Stanislas Larnier

Independent research

Guillaume Cassez — ORCID 0009-0007-0987-3931 · cassez.guillaume@gmail.com · guillaume-cassez.fr

Stanislas Larnier — stanislaslarnier@gmail.com · HAL stanislas-larnier

Author list established 2026-10-04 — Stanislas Larnier joins this paper in second position, by mutual agreement between the two authors, as on the four BRATS papers of the same programme. Versions deposited on Zenodo before that date carry Guillaume Cassez alone. No number in the manuscript is modified.

Cityscapes val · ConvNeXt-V2-Base + UPerNet · CE + Dice + 0.5·blob (Kofler, IPMI 2023) vs CE + Dice · 3 seeds × 160 epochs at 1024×2048

Abstract

We report a pre-registered, controlled evaluation of the **blob loss** of Kofler *et al.* [IPMI 2023] — an auxiliary term that gives every ground-truth instance the same weight regardless of its pixel count — ported to **exclusive multi-class semantic segmentation at the native Cityscapes resolution** (1024×2048, 19 classes, softmax regime). Arm **G** (CE + Dice + 0.5·blob) is compared against its exact paired reference **B** (CE + Dice): same ConvNeXt-V2-Base + UPerNet architecture, same 160-epoch recipe, same augmentation distribution, three shared seeds (42, 123, 456) — only the loss changes.

The **pre-registered primary endpoint is null**: dataset-level official mIoU on the shared 500-image holdout, $\Delta(G-B) = +0.170$ pt, 95 % CI [-0.233 ; +0.551], two-sided paired image-bootstrap $p = 0.4032$ ($B = 10\,000$ replicates). Equal weight per instance does **not** improve global mIoU in this regime. What the term actually does is a **measured trade-off**. It *buys* thin-object pixel coverage — traffic light +0.98, pole +0.58, bicycle +0.57 IoU vs B (Holm-significant within the 19-class family), truck +3.78 (CI excluding zero), Boundary F1 (3 px) +0.63 — and it *pays* in pedestrian instance integrity: strict pedestrian recall **-4.34**, small-instance stratum (T1) **-5.65**, crowd-group instances **-5.74**, pedestrian pixel precision **-2.70** (all vs paired B, Holm = 0), and a general **fragmentation** of the masks: connected components rise from 615.1 to 933.9 per image (×1.52), up in **16 of 19 classes** against the paired arm B (14 Holm-significant rises; the only material decrease is terrain, -1.8) and in **19 of 19** against the control arm (person masks ×2.2) — a new per-class decomposition that *revises* the working mechanistic hypothesis of the program (the term does not prune small components; it creates holes). On the program's pre-registered versatility criterion (36 endpoints × 13 arms), G is a **dominated specialist**: mean percentile 41.9, worst rank 13/13, maximal damage -6.76 pt, 13 significant losses against 4 significant gains.

The same verdict holds on a second dataset and a second probability regime: on BRATS 2023 (region-based sigmoid, MedNeXt, nnU-Net recipe, 5-fold CV, $n = 1\,196$), the blob arm alone scored **-0.00376 Dice** vs baseline, yet as **expert 3 of the winning MoE-V3** it contributed **+0.00566 Dice, first of 29 arms** [DOI 10.5281/zenodo.22903668]. On the Cityscapes program plateau, the expert-mixture arm initialised from these specialists (MoE-V3-CS, master table P3.14) points the same way, as context: $\Delta mIoU +0.45$ pt [+0.11 ; +0.82], $p = 0.0066$ vs its recipe-paired control, maximal damage -0.53 pt, small-instance recall T1 +0.80, traffic light not degraded — the profile of an arm that puts G's thin-object speciality to use without carrying its damage. The dedicated analysis of that mixture arm (routing diagnostics, multiplicity families) is not the subject of this paper.

Contributions. (1) A pre-registered **null primary** for equal-weight-per-instance auxiliary loss in full-resolution exclusive-softmax segmentation, reported as such. (2) A complete **trade-off diagnostic**: per-class IoU, seven business-metric families on pedestrian instances, and a **new per-class connected-component decomposition** (933.9 vs 615.1 components/image) that contradicts the initially hypothesised compaction mechanism — the paper reports what the tables show. (3) An **exact port** of Kofler's algebra (eq. 1) to the exclusive-softmax regime with precomputed instance packs: +1–2 % epoch time against +870 s/epoch for the naive path, naive↔precomputed parity locked by unit tests, and the horizontal-flip alignment pitfall documented. (4) Evidence, on two datasets and two probability regimes, that the term **pays as an initialised expert of a mixture** while null-or-harmful alone. (5) Public release of code, configs, tables, figures and regeneration scripts.

1. Introduction

Voxel-wise losses — cross-entropy, Dice — weight every ground-truth instance by its volume. Small instances (rare classes, distant or partially occluded objects) contribute few pixels to the loss and are, in a strict optimisation sense, *rationaly sacrificed*. Kofler *et al.* [2023] proposed the **blob loss**: an auxiliary term that computes a soft binary dice **per ground-truth instance**, over a domain restricted to that instance plus all pixels outside other instances of the same class, and averages instances $\rightarrow$ class $\rightarrow$ image $\rightarrow$ batch. Every instance then carries the same total weight regardless of its pixel count. The term was validated in biomedical segmentation under the nnU-Net **region-based** recipe (overlapping WT/TC/ET targets, sigmoid channels), where it improved small-structure metrics.

This paper asks the transfer question on the other side of the fence:

Does equal weight per instance pay in **exclusive multi-class** semantic segmentation — softmax over 19 flat classes — at the **native Cityscapes resolution**, under a fixed training budget and a pre-registered primary endpoint?

The setting is a controlled arm of a four-paper loss program on Cityscapes at 1024×2048 with ConvNeXt-V2-Base + UPerNet: a distance-map auxiliary regression study [Cassez 2026a, DOI 10.5281/zenodo.21006236], a boundary-loss ablation whose arm B (CE + Dice, 160 epochs, seeds 42/123/456) is the exact paired reference used here [Cassez 2026b, DOI 10.5281/zenodo.21006393], this blob-loss study, and a mixture-of-experts arm (MoE-V3-CS) initialised from these specialists, of which the present arm G is expert 3 and whose results enter here only as plateau context (§6.1, §6.5). The program mirrors a completed BRATS track [DOI 10.5281/zenodo.22903668; 10.5281/zenodo.22906447; 10.5281/zenodo.22904810; 10.5281/zenodo.20110976], so every blob arm exists on both datasets under two probability regimes.

The answer, measured before this manuscript was written and pre-registered in the program’s artifact chain, is **no** — **the primary endpoint is null** (Δ mIoU +0.170 pt, $p = 0.4032$). This is therefore a negative-result paper in the house discipline of the program: negative and neutral results are reported as such, all official metrics are shown, and the contribution is the **diagnostic** (what the term buys, what it breaks, and the mechanism the tables actually support), the **exact port** (algebra, cost, pitfalls), and the **usage** (an initialised expert of a mixture, measured on two datasets).

2. Related work

Class imbalance in semantic segmentation. Cross-entropy with class re-weighting is the standard remedy; we use ISNS weights (inverse square-root of the number of samples) in *both* arms, so the comparison isolates the blob term on top of an already re-weighted CE. Dice loss [Milletari 2016] optimises region overlap and is class-balanced at the *class* level, not at the *instance* level: within a class, one large instance still dominates many small ones. Focal loss [Lin 2017] re-weights hard pixels, again pixel-wise. The blob loss is the imbalance remedy that operates **per instance**.

Blob loss. Kofler *et al.* [2023, IPMI; arXiv:2205.08209] introduce the term for biomedical segmentation under the nnU-Net region-based recipe [Isensee 2021], with total loss $\mathcal{L} = \mathcal{L}_{seg} + \beta \mathcal{L}_{blob}$ and the paper’s ratio $\alpha : \beta = 2 : 1$. Their validation regime is **sigmoid** (independent per-region channels, overlapping regions). The exclusive-softmax regime studied here puts classes in direct competition for every pixel; §7 discusses why the verdict differs from the intuition the term invites.

Boundary-aware and distance-based companions. Within the same program and hardware: the Kervadec boundary loss [Kervadec 2019] ablation [Cassez 2026b] and the distance-map auxiliary regression [Cassez 2026a] both report their pre-registered endpoints on the identical holdout and bootstrap protocol, which makes the present arm directly comparable to them (13-arm context table, §6.1).

Fragmentation and consensus filtering. Connected-component counts are a spatial-coherence proxy to which dataset-level mIoU is blind. The companion boundary paper measures a CC **consensus veto** ($C\oslash B$) that prunes −18.7 % of spurious fragments at no mIoU cost [Cassez 2026b]; on BRATS the analogous rule is published in [DOI 10.5281/zenodo.22904810]. The per-class fragment decomposition introduced here (§6.4) is, to our knowledge, the first measurement of how an instance-balancing loss moves fragmentation *per class* on Cityscapes.

Mixture-of-experts from initialised specialists. The BRATS MoE-V3 [DOI 10.5281/zenodo.22903668] trains a patch-wise gate over four experts initialised from independently trained specialist arms — including the blob arm. The Cityscapes program plateau transposes the same design (four experts B, D, Dp, G; top-2 gate over 3×3 patches; 80 epochs from expert weights): it is the MoE-V3-CS arm of the master tables P3.14/P3.17, of which the present arm G is expert 3. This paper cites the published BRATS track for the *usage* half of its thesis and uses the MoE-V3-CS arm only as plateau context; the dedicated analysis of that arm (routing diagnostics, primary endpoint against its control) is not its subject.

3. Method — the blob loss in the exclusive-softmax regime

Let Ω be the image domain, $C = 19$ the number of classes, $p_c(x) \in [0, 1]$ the **softmax** probability of class c at pixel x (so $\sum_c p_c(x) = 1$ — the regime difference with BRATS), and $y(x) \in \{0, \dots, 18, 255\}$ the ground-truth train id (255 = void/ignore). For class c , the ground-truth mask splits into connected **instances** $n \in I_c$ (8-connectivity on the label map; group annotations of the Cityscapes instance maps are handled at the business-metric level, §5.3, not inside the loss).

Per-instance domain. Following Kofler *et al.* eq. 1 exactly, the domain of instance n of class c is the whole image *minus the other instances of the same class*:

$$\Omega_n = (\{x : y(x) \neq 255\} \setminus \bigcup_{m \in I_c, m \neq n} m) \cup n.$$

Per-instance soft binary dice. With t_n the indicator of instance n , $s_n = \sum_{x \in n} p_c(x)$ the predicted mass on the instance, $c_n = |n|$ its pixel count, and $S_0^{(c)} = \sum_{x \in \Omega_n \setminus n} p_c(x)$ the mass on the instance-free part of its domain (so $\sum_{x \in \Omega_n} p_c = S_0^{(c)} + s_n$):

$$\ell_n = 1 - \frac{2s_n + \varepsilon}{S_0^{(c)} + s_n + c_n + \varepsilon}, \quad \varepsilon = 1.0 \text{ (nnU-Net convention)}.$$

Ignored pixels (255) fall in a trash bucket: they are excluded from both $S_0^{(c)}$ and the instances, and are never counted as false negatives (the valid semantics of the BRATS port).

Two-level mean. Faithful to the BRATS RegionBlobLoss the port comes from:

$$\mathcal{L}_{blob} = \frac{1}{|\mathcal{B}|} \sum_{img \in \mathcal{B}} \frac{1}{|C_{img}|} \sum_{c \in C_{img}} \frac{1}{|I_c|} \sum_{n \in I_c} \ell_n,$$

where C_{img} is the set of classes present in the image. Every instance of every present class carries the same total weight per image — the defining property of the term. `min_blob_pixels = 0`: no instance is filtered by size, since small instances are precisely what the term is meant to rescue.

Composite losses (arms G and B). With CE at ISNS class weights and Dice (smooth 1.0), both ignoring 255:

$$\mathcal{L}_G = 0.5 \mathcal{L}_{CE}^{ISNS} + 0.5 \mathcal{L}_{Dice} + 0.5 \mathcal{L}_{blob}, \quad \mathcal{L}_B = 0.5 \mathcal{L}_{CE}^{ISNS} + 0.5 \mathcal{L}_{Dice}.$$

$\beta = 0.5$ is Kofler’s default ($\alpha : \beta = 2 : 1$ with the segmentation loss weighing 1.0 in total). **No sweep of β is performed** — the fixed-budget comparison stance of the whole program (same position as the companion papers on λ_b and the distance-map weight).

4. Implementation — precomputed instance packs

Two execution paths, same algebra. (1) *Precomputed* (used for training arm G): per-image instance labels are packed offline into `.blob.npz` files (`blob_glob/blob_counts/blob_csr`, built by `src/losses/blob_lab.py` via `scripts/precompute_blob_lab.py`) and delivered by the dataloader; the forward pass is then elementary GPU scatter-add kernels — the CPU connected-component pass and the 26 small per-class GPU synchronisations disappear from the training loop. (2) *On-the-fly fallback*: without a pack, the same packing functions recompute it from the target — this is the correctness reference used by tests and evaluation harnesses, giving structural parity between the two paths.

Measured cost (RTX PRO 6000 Blackwell, full-resolution ground truth, 2026-09-17): the naive on-the-fly path costs **+870 s/epoch** — unacceptable; the precomputed path costs **+1–2 %/epoch**. The connected-component labelling itself runs at **22.4 ms/image** with `cc3d` (8-connectivity; 16 classes present, 118 instances on a representative full-resolution image) against 59.7 ms with `scipy.ndimage.label`. Packs are computed once for the 2 975-image train split (labels are shared across seeds).

Parity and regeneration. Naive↔precomputed parity is locked by `tests/test_blob_loss.py`. All 2 975 packs were regenerated **bit-identical** on 2026-09-26 (verified over the full set). The packs themselves are heavy and are not shipped in the release bundle; `scripts/precompute_blob_lab.py` regenerates them identically.

The flip pitfall (documented). The instance pack must undergo the **same** horizontal flip as image and label. In arm G the dataset applies the flip itself ($p = 0.5$, jointly on image + label + pack) and the albumentations flip is **disabled** in the config, so the blob glob can never be misaligned. The augmentation *distribution* is identical to arm B’s (hflip $p = 0.5$, photometric jitter, Gaussian blur); only the application point differs. A sanity inherited from the program’s SDT audit: all three resolved G configs have the structural flip OFF — the arm is unaffected by the misalignment bug found in earlier boundary/distmap arms.

5. Experimental setup

5.1 Model, recipe, budget

Architecture (identical in both arms): ConvNeXt-V2-Base [Woo 2023] pretrained ImageNet-22K (FCMAE) then fine-tuned on ImageNet-1K, + UPerNet head [Xiao 2018] with auxiliary FCN deep supervision (weight 0.4), 19 logits per pixel at full input resolution. **Training:** 160 epochs, batch $2 \times$ gradient accumulation 4 (effective 8), AdamW (lr 6×10^{-5} , weight decay 0.01), polynomial decay (power 1.0, 1-epoch warmup), BF16 autocast, channels_last, inputs at native 1024×2048 without cropping. Seeds **42, 123, 456** in both arms. Arm G trained 17–20 September 2026.

Measured cost of arm G (per seed, results/p3_queue/pilot_fullres_G_blob_seed{42,123,456}.log): 160/160 epochs, median **656 s/epoch** [655–669], VRAM **31.5 GB**, final train loss 0.3743 / 0.3735 / 0.3765, $\approx$ **29.2 h per seed**, $\approx$ **87 h GPU** for the arm.

5.2 Data and holdout

Cityscapes fine annotations [Cordts 2016]: 2 975 train / 500 val images, 19 evaluation classes. All evaluations run on the program’s **pre-specified shared holdout**: the first 500 val images in loader order (first:500) — identical for all 13 arms of the program, fixed before any arm was trained. The official Cityscapes leaderboard was **not** submitted to (declared limitation, §8).

5.3 Metrics

- **mIoU** — dataset-level, from a single aggregated confusion matrix, estimator validated **bit-identical to the official cityscapesScripts** routine (tests/test_official_miou.py). Per-class IoU from the same matrices.
- **Boundary F1 (3 px)** — per-class F1 of predicted vs ground-truth contours within a 3-pixel tolerance [protocol of Perazzi 2016], averaged over classes present per image; also restricted to {person, rider}.
- **Pedestrian instance metrics** — from the official *_gtFine_instanceIds.png maps, classes {person, rider}: pixel recall *any-ped* (prediction $\in$ {person, rider}) and *strict* (prediction = the instance’s own class); detection rate at recall ≥ 0.5 ; pedestrian **pixel precision** (fraction of predicted-pedestrian pixels falling inside a ground-truth pedestrian instance).
- **Strata** — all instances / individual instances (instId ≥ 1000) / **crowd-group** instances (instId < 1000 — the official ambiguous/tightly-packed pedestrian groups, the best available proxy for partially occluded pedestrians) / **size terciles** of individual instances, T1 = smallest tercile (proxy for distance and occlusion), T3 = largest.
- **Fragments** — 8-connected components per image summed over the 19 class masks (count_fragments, src/postprocessing/consensus.py), plus a **new per-class decomposition** introduced by this paper (components per class per image, arms G/B/control).

5.4 Pre-registered primary endpoint and statistics

The single pre-registered primary endpoint is the **dataset-level official mIoU, arm G vs its paired reference B** (paired = same architecture, recipe, budget, augmentation, seeds — only the loss differs). Test: **paired image-bootstrap** over the 500 holdout images — positions resampled with replacement, **B = 10 000** replicates, bootstrap seed **20260618**, the three seeds averaged *within each replicate*, 95 % percentile CI, two-sided $p = 2 \cdot \min(\text{frac } \Delta \leq 0, \text{frac } \Delta \geq 0)$. The primary family is the single pair, so Holm = p .

Secondary/exploratory families, each with its Holm correction declared where reported: per-class IoU (19-class family per arm pair), business metrics (Holm(2) over this paper’s two pairs G-vs-B and G-vs-control), per-class fragmentation (19-class family), 13-arm mIoU context (**two** families, both given side by side in §6.1: the 12-pair family of the master table results/moe_v3_cs/p314/master_table.json and the program’s 15-pair exploratory family of results/moe_v3_cs/metiers_experts/table_metiers_experts.json).

Two significance criteria, never conflated. *Holm-significant* (corrected $p < 0.05$ within the stated family) and *CI-excluding-zero* (raw 95 % CI) are reported **separately** everywhere — bold in tables marks Holm significance only, and the p / Holm columns carry the exact values. This distinction matters for this arm: e.g. truck +3.78 IoU vs B excludes zero but does not survive the 19-class Holm (0.460), while traffic light +0.98 survives (Holm < 0.001).

The two comparison pairs. The budget-paired reference is **B** (160 epochs) — the pair of the primary endpoint. The 13-arm program context uses a different reference, the MoE **control** (CE + Dice, 80 epochs, B-initialised) — *not* budget-paired with G, always labelled as context, never substituted for the paired pair. Both are reported side by side throughout.

Non-determinism, declared. cuDNN benchmark is on (deterministic: false, BF16). The primary numbers come from the harness forward of the pre-registered artifact; the business-metric regeneration (§6.3) comes from a second independent GPU forward of the *same* checkpoints. The two agree within ≤ 0.01 pt on every shared quantity (per-seed mIoU gaps $\leq 2 \times 10^{-4}$ pt; primary $\Delta +0.170$ vs $+0.173$, p 0.4032 vs 0.3952 — same verdict), of the same order as the 4.1×10^{-3} pt already documented in the program’s consolidation sanity.

6. Results

6.1 Primary endpoint: null

Pre-registered primary, paired image-bootstrap over the 500-image holdout (artifact `results/moe_v3_cs/harness/table_Gseul_P310.json`, 2026-09-22):

Quantity	G — CE+Dice+0.5·blob	B — CE+Dice
mIoU (point)	81.263	81.093
95 % CI	[79.815 ; 82.431]	[79.694 ; 82.204]
mIoU seed 42	81.297	81.429
mIoU seed 123	81.291	80.861
mIoU seed 456	81.202	80.989

$\Delta(G-B) = +0.170$ pt · 95 % CI $[-0.233 ; +0.551]$ · $p = 0.4032$ (single pair, Holm = p) → **not significant**.

The honest reading: the primary endpoint is **null**. The blob loss does not improve mIoU in full-resolution exclusive-softmax semantic segmentation. Per-seed values do not hide a consistent direction either: G wins seed 123 (+0.43) and seed 456 (+0.21), loses seed 42 (−0.13).

The second, independent GPU forward of the same checkpoints (business-metric regeneration path, §6.3) gives $\Delta = +0.173$ pt $[-0.229 ; +0.554]$, $p = 0.3952$ — same verdict, gap ≤ 0.01 pt (declared cuDNN/BF16 non-determinism).

13-arm program context (reference: recipe-paired control of the mixture arm, CE + Dice 80 epochs — *not* budget-paired with G; artifact `results/moe_v3_cs/p314/master_table.json`):

#	Arm	mIoU	Δ vs control	p	Holm (12-pair family)	Holm (15-pair family)
1	D · CE+Kervadec EDT	81.69	+0.52 [+0.14 ; +0.91]	0.0048	0.058	0.075
2	consensus D $\oslash$ B	81.65	+0.48 [+0.11 ; +0.87]	0.0082	0.082	0.107
3	consensus Dp $\oslash$ B	81.65	+0.48 [+0.02 ; +0.96]	0.0396	0.317	0.414
4	Dp · CE+SDT (distmap)	81.64	+0.47 [+0.04 ; +0.93]	0.0312	0.281	0.396
5	MoE-V3-CS (4 experts)	81.62	+0.45 [+0.11 ; +0.82]	0.0066	0.073	0.092
6	A · CE alone	81.28	+0.11 [−0.27 ; +0.52]	0.5722	1.000	n.a. (a)
7	G · CE+Dice+Blob (this paper)	81.26	+0.10 [−0.28 ; +0.47]	0.6552	1.000	1.000
8	consensus C $\oslash$ B	81.24	+0.07 [−0.34 ; +0.47]	0.7758	1.000	1.000
9	C · CE+Dice+Kervadec EDT	81.23	+0.06 [−0.34 ; +0.47]	0.8096	1.000	1.000
10	B · CE+Dice (160 ep)	81.09	−0.08 [−0.45 ; +0.28]	0.6748	1.000	1.000
11	consensus Cp $\oslash$ B	81.01	−0.15 [−0.52 ; +0.17]	0.3498	1.000	1.000
12	Cp · CE+Dice+SDT	80.89	−0.28 [−0.60 ; +0.02]	0.0628	0.440	0.618
—	control (ref, CE+Dice 80 ep)	81.17	—	—	—	—

(a) A is outside the 15-pair family (partial coverage; `annexe_couverture_partielle` of P3.17 = [‘A’]), so no Holm value exists for it there.

Two distinct multiplicity families, both reported. The *12-pair* family is the master table itself (each of the 12 arms against the control, `results/moe_v3_cs/p314/master_table.json`). The *15-pair* family is the program’s exploratory family, which adds

the 4 fusion-versus-expert comparisons (results/moe_v3_cs/metiers_experts/table_metiers_experts.json, carried over by P3.17). Family sizes are computed from the artifacts and each Holm value is recomputed from the raw p-values and compared against the stored one (blocking check, papers/paper4/tables/SANITY.md).

G ranks **7th of 13** by amplitude (the control, $\Delta = 0$, ranks 10th). **No arm passes Holm 0.05 in either family** — best 0.0576 over 12 pairs, 0.0750 over 15 pairs — stated as-is: on mIoU, the conclusions of this program rest on **amplitudes and their CIs**, not on post-Holm significance.

6.2 Per-class IoU: thin-object coverage bought




Figure 1: F1 — Per-class forest plot, G vs B paired

Figure 1: per-class Δ IoU, arm G vs paired reference B (bootstrap apparié, $B = 10\,000$). Filled markers = Holm-significant within the 19-class family; open markers = CI excluding zero without Holm survival; the two criteria are plotted separately, never merged.

Full 19-class table (paired G–B; context G–control from the same artifact family). Bold = Holm (19-class family) < 0.05 :

Class	IoU G	IoU B	$\Delta(G-B)$	95 % CI	p	Holm	$\Delta(G-ctrl)$	95 % CI	p	Holm
road	98.45	98.44	+0.01	[−0.05 ; +0.06]	0.7970	1.000	+0.04	[−0.02 ; +0.10]	0.1874	1.000
sidewalk	87.04	87.07	−0.03	[−0.44 ; +0.38]	0.9184	1.000	+0.06	[−0.31 ; +0.48]	0.7642	1.000
building	93.32	93.27	+0.05	[−0.13 ; +0.18]	0.5170	1.000	+0.07	[−0.03 ; +0.18]	0.1612	1.000
wall	57.21	57.92	−0.70	[−4.97 ; +2.43]	0.7974	1.000	+1.07	[−1.32 ; +4.05]	0.4916	1.000
fence	65.35	66.07	−0.72	[−2.47 ; +0.71]	0.3874	1.000	−0.46	[−1.29 ; +0.34]	0.2668	1.000
pole	70.34	69.76	+0.58	[+0.34 ; +0.81]	<0.0001	<0.001	+0.29	[+0.06 ; +0.51]	0.0136	0.231
traffic light	77.83	76.85	+0.98	[+0.58 ; +1.37]	<0.0001	<0.001	+0.47	[+0.12 ; +0.82]	0.0108	0.194
traffic sign	83.92	83.52	+0.40	[−0.04 ; +0.83]	0.0686	0.706	+0.21	[−0.13 ; +0.54]	0.2182	1.000
vegetation	93.15	93.09	+0.06	[−0.00 ; +0.13]	0.0588	0.706	+0.02	[−0.03 ; +0.07]	0.4260	1.000
terrain	66.40	66.10	+0.30	[−0.46 ; +1.06]	0.4454	1.000	+0.01	[−0.61 ; +0.60]	0.9850	1.000
sky	95.64	95.60	+0.05	[−0.05 ; +0.15]	0.3538	1.000	+0.11	[+0.00 ; +0.22]	0.0410	0.656
person	85.22	84.90	+0.32	[+0.07 ; +0.56]	0.0146	0.234	+0.07	[−0.18 ; +0.30]	0.5824	1.000
rider	67.58	68.49	−0.90	[−1.98 ; +0.03]	0.0632	0.706	−1.34	[−2.42 ; −0.35]	0.0072	0.137
car	95.71	95.46	+0.25	[+0.03 ; +0.54]	0.0212	0.297	+0.14	[−0.13 ; +0.45]	0.3136	1.000
truck	84.03	80.24	+3.78	[+0.09 ; +8.30]	0.0354	0.460	+3.43	[−1.22 ; +8.31]	0.1826	1.000
bus	90.07	89.75	+0.32	[−1.38 ; +2.20]	0.7154	1.000	−1.48	[−4.74 ; +1.01]	0.3274	1.000
train	81.20	83.74	−2.54	[−5.59 ; −0.34]	0.0176	0.264	−1.04	[−2.79 ; +0.86]	0.2966	1.000
motorcycle	71.15	70.70	+0.46	[−0.90 ; +2.31]	0.5112	1.000	+0.07	[−0.93 ; +0.99]	0.9362	1.000
bicycle	80.38	79.81	+0.57	[+0.28 ; +0.90]	<0.0001	<0.001	+0.05	[−0.17 ; +0.29]	0.6634	1.000

Against paired **B**: the Holm survivors are exactly the thin, signal-rich classes — **traffic light +0.98**, **pole +0.58**, **bicycle +0.57**. Truck (+3.78), person (+0.32) and car (+0.25) exclude zero but do not survive the 19-class Holm (0.460 / 0.234 / 0.297); train (−2.54) is the single degradation excluding zero (Holm 0.264). Against the **control**, G is the **only arm of the 13** that improves traffic light (+0.47, CI excluding zero; the boundary/distmap arms all *degrade* it by −2.3 to −2.8), at the price of rider (−1.34, CI excluding zero; Holm(2) 0.0136 in §6.3).

6.3 Business metrics: pedestrian integrity paid

F2 — The trade-off of the blob loss alone: pixel IoU of fine objects $\uparrow$ against pedestrian instance recall $\downarrow$ — and the topology bill: more connected components in 16 of 19 classes vs paired B (19 of 19 vs control, pedestrians $\times 2.2$) G vs paired B (160 epochs, only the loss changes); left: per-class attribution P3.16 · middle: regenerated instance metrics · right: per-class decomposition of connected components (novel measure, P3.16 preds)




Figure 2: F2 — The measured trade-off

Figure 2: the trade-off, both pairs. Left/middle: thin-object pixel coverage (Boundary F1, per-class IoU) goes up; right: pedestrian instance recall (strict, size strata, crowd) and pedestrian pixel precision go down. CIs from the paired image-bootstrap.

Regenerated from the per-image npz artifacts (500 images $\times$ 3 seeds per arm; bit-identical to the program's cross-arm attribution artifacts on 19 shared metrics; ≤ 0.01 pt vs the primary harness forward on the shared mIoU row — 14/14 blocking sanity checks, papers/paper4/tables/SANITY.md). Units: points ($\times 100$), except fragments = components per image. **Bold** = Holm(2) < 0.05 over this paper's two pairs; n = images contributing to the metric:

Metric	n	G	B	ctrl	$\Delta(G-B)$ [95 % CI]	p	Holm(2)	$\Delta(G-ctrl)$ [95 % CI]	p	Holm(2)
mIoU	500	81.26	81.09	81.17	+0.173 [-0.229 ; +0.554]	0.3952	0.7904	+0.094 [-0.281 ; +0.466]	0.6592	0.7904
IoU person	500	85.22	84.88	85.15	+0.344 [+0.092 ; +0.578]	0.0088	0.0176	+0.067 [-0.182 ; +0.302]	0.5838	0.5838
IoU rider	500	67.58	68.49	68.92	-0.911 [-2.003 ; +0.036]	0.0612	0.0612	-1.339 [-2.413 ; -0.352]	0.0068	0.0136
Boundary F1 3 px	500	77.00	76.37	76.63	+0.630 [+0.454 ; +0.806]	<0.0001	<0.001	+0.376 [+0.216 ; +0.537]	<0.0001	<0.001
Boundary F1 3 px, ped	500	67.73	67.30	67.49	+0.431 [-0.119 ; +0.953]	0.1204	0.2408	+0.247 [-0.201 ; +0.702]	0.2814	0.2814
Ped. strict recall	441	72.84	77.17	76.25	-4.335 [-4.948 ; -3.729]	<0.0001	<0.001	-3.417 [-3.967 ; -2.868]	<0.0001	<0.001
Ped. pixel precision	473	63.27	65.96	70.03	-2.699 [-3.484 ; -1.951]	<0.0001	<0.001	-6.763 [-7.800 ; -5.785]	<0.0001	<0.001
All-instances recall	441	76.55	80.99	79.98	-4.439 [-5.000 ; -3.884]	<0.0001	<0.001	-3.423 [-3.911 ; -2.942]	<0.0001	<0.001
All-instances det@0.5	441	85.12	88.05	87.11	-2.936 [-3.817 ; -2.122]	<0.0001	<0.001	-1.994 [-2.696 ; -1.328]	<0.0001	<0.001
Individual recall	440	76.74	81.11	80.10	-4.376 [-4.935 ; -3.823]	<0.0001	<0.001	-3.361 [-3.853 ; -2.866]	<0.0001	<0.001
Individual det@0.5	440	85.17	88.08	87.18	-2.903 [-3.819 ; -2.069]	<0.0001	<0.001	-2.002 [-2.712 ; -1.302]	<0.0001	<0.001
Crowd recall	63	69.97	75.71	74.58	-5.736 [-7.864 ; -4.033]	<0.0001	<0.001	-4.612 [-6.531 ; -2.940]	<0.0001	<0.001
Crowd det@0.5	63	82.01	86.24	83.60	-4.233 [-7.937 ; -1.058]	0.0046	0.0092	-1.587 [-4.233 ; +0.529]	0.2250	0.2250
T1 (smallest) recall	315	60.91	66.56	64.99	-5.645 [-6.650 ; -4.627]	<0.0001	<0.001	-4.080 [-5.092 ; -3.091]	<0.0001	<0.001
T1 det@0.5	315	68.26	72.84	71.73	-4.574 [-6.218 ; -2.890]	<0.0001	<0.001	-3.466 [-5.017 ; -1.960]	<0.0001	<0.001
T2 recall	320	80.88	85.46	84.53	-4.584 [-5.316 ; -3.880]	<0.0001	<0.001	-3.656 [-4.370 ; -2.938]	<0.0001	<0.001
T2 det@0.5	320	91.08	93.78	93.15	-2.695 [-3.893 ; -1.609]	<0.0001	<0.001	-2.064 [-3.196 ; -0.928]	0.0004	0.0004
T3 (largest) recall	335	91.68	93.85	93.47	-2.175 [-2.478 ; -1.875]	<0.0001	<0.001	-1.798 [-2.090 ; -1.510]	<0.0001	<0.001
T3 det@0.5	335	98.78	98.96	98.96	-0.180 [-0.697 ; +0.238]	0.4654	0.4654	-0.177 [-0.468 ; +0.054]	0.1586	0.3172
Fragments (comp./img)	500	933.88	615.13	519.26	+318.8 [+305.4 ; +331.9]	<0.0001	<0.001	+414.6 [+397.8 ; +431.1]	<0.0001	<0.001

The trade-off in one sentence: the term **buys** contour quality (Boundary F1 +0.63 vs B, Holm-significant) and thin-object pixel coverage (§6.2), and **pays** in every pedestrian-instance integrity metric — strict recall -4.34, smallest tercile -5.65, crowd groups -5.74, pixel precision -2.70, detection@0.5 -2.94 (all vs paired B, Holm(2) < 0.001 except crowd det@0.5 at 0.0092). The payment is ordered by instance size: T1 -5.65 > T2 -4.58 > T3 -2.18 — **the smallest instances lose the most coverage**, the exact opposite of the term’s intent. Person IoU does rise (+0.344, Holm(2)-significant vs B): the pixel *mass* of the person class grows while the *integrity* of individual pedestrian masks degrades — IoU and instance recall move in opposite directions, which is precisely why the mIoU-centric reading alone would have missed the effect.

6.4 Fragmentation: the mechanism the tables show

New per-class measurement of this paper (500 images × 3 seeds, arms G / B / control, 8-connected components per class per image, paired bootstrap + 19-class Holm; full table in Appendix B):

Reading	Value (G vs B)
Total components per image	933.9 vs 615.1 — ×1.52 (+318.8 [+305.4 ; +331.9], Holm < 0.001)
Classes with more components	16 of 19 vs paired B (14 Holm-significant rises); 19 of 19 vs the control arm. The three that do not rise vs B: terrain -1.8 (Holm < 0.001, the only material decrease), wall -0.02 (p = 0.95) and bus -0.02 (p = 0.70), both null
Largest absolute rises	car +49.4, pole +45.4, vegetation +43.8, road +38.6, traffic sign +38.4, building +32.5, sidewalk +30.8
Pedestrian masks	person 24.3 → 53.2 components/image (×2.2); rider 2.0 → 4.1
Thin classes that <i>gain</i> IoU	also gain components: traffic light +0.9, bicycle +2.2 (Holm-significant)

Revision note (house discipline: the paper reports what the tables show). The program’s initial working hypothesis — inherited from the BRATS intuition that an instance-balancing term suppresses small spurious components in favour of compact thin objects — is **contradicted** by this decomposition: components *increase* almost everywhere, including in the classes whose pixel IoU improves. The mechanism the measurements support is the opposite of compaction: in the exclusive-softmax regime, the equal-weight-per-instance gradient buys **dispersed pixel coverage** for thin and ample classes (per-class IoU ↑, Boundary F1 ↑) at the price of **holes inside ground-truth instances** — a holed mask splits into disjoint components (fragments ×1.52), which mechanically lowers instance-level recall (strict -4.34, T1 -5.65) and pedestrian pixel precision (-2.70). No gradient-level analysis is performed here; §8 lists this as a limitation.

6.5 Versatility ranking: a dominated specialist

The program’s pre-registered versatility criterion: 36 endpoints (19 per-class IoU + 17 business metrics) $\times$ 13 arms, same holdout and bootstrap; “maximal damage” = the worst endpoint-level Δ vs control (closer to 0 = no weak point). G’s row and its neighbours (full ranking, artifact results/moe_v3_cs/p317_polyvalence/table_polyvalence.json):

Rank	Arm	Mean percentile	Worst rank	#1	Max damage (pt)	z	Max peak (pt)	Δ IoU (p / Holm)	Gains sig. (raw/Holm)	Losses sig. (raw/Holm)	Pareto
1	MoE-V3-CS	60.4	11	0	-0.53 (crowd det@0.5)	-0.35	+4.24 (truck IoU)	+0.45 (0.0066 / 0.092)	7/1	2/1	non-dominated
5	B · CE+Dice	51.2	13	2	-4.06 (ped precision)	-0.63	+2.65 (crowd det@0.5)	-0.08 (0.6592 / 1.000)	11/8	6/6	dominated
8	G · CE+Dice+Blob (this paper)	41.9	13	7	-6.76 (ped precision)	-1.05	+3.43 (truck IoU)	+0.09 (0.6592 / 1.000)	4/1	13/12	dominated
11	D · CE+Kervadec EDT	77.3	13	16	-22.20 (ped precision)	-3.44	+5.49 (wall IoU)	+0.52 (0.0050 / 0.075)	17/13	4/3	non-dominated

G’s profile: mean percentile **41.9** (mean rank 8.2), worst rank **13/13**, last on 16 of 36 endpoints — yet **#1 on 7 of 36 endpoints**, more than any arm except D (16): a genuine specialist. The damage is concentrated on the pedestrian family (worst endpoint: pedestrian pixel precision **-6.76 pt**, $z = -1.05$ inter-arm σ ; worst family mean: pedestrian -5.09; best family: contours +0.31). The **price per mIoU point** — maximal damage divided by Δ mIoU — is **-71.7 pt** for G against **-1.2 pt** for the MoE-V3-CS mixture arm: at a comparable global gain, the mixture does not carry the hole. G is **Pareto-dominated** (by B, C, C $\oslash$ B, Dp $\oslash$ B, the control and MoE-V3-CS). The pre-registered T0 verdict (significant Δ mIoU *and* no endpoint > 1 pt below reference *and* no class degraded > 0.5 pt) is reported **as computed: True for the mixture arm, and G is its antithesis** — a dominated specialist whose speciality (thin objects) is precisely what a mixture initialised from these arms puts to use (§7.2).

F3 — Position of G in the P3.17 versatility ranking (13 arms, 36 endpoints)
 G = dominated specialist (mean percentile 41.9 · worst rank 13/13 · max damage -6.76 pt on precision_ped_pixels);
 the only arm in the T0 zone is the expert-initialised mixture built from these arms (MoE-V3-CS, §7.2)




Figure 3: F3 — G in the versatility ranking

Figure 3: position of arm G in the 36-endpoint $\times$ 13-arm ranking, against the mixture MoE-V3-CS and the plateau. Specialist peak (#1 on 7 endpoints) and maximal damage (-6.76 pt) on the same arm.

6.6 Cross-dataset: the same verdict on BRATS

Numbers **cited** from the published BRATS program (second dataset, second probability regime — region-based sigmoid, MedNeXt, nnU-Net recipe, 5-fold CV over $n = 1\,196$; sources verified 2026-09-30, not re-measured here):

BRATS arm	Result	Source
blob alone vs nnU-Net baseline	-0.00376 Dice (5-fold CV; 0/5 folds won; last of the 4 expert arms)	BRATS papers/paper3/versions/v3/NOTE_DECISION_V3.md l.46 + tables/v3_vs_consensus.md l.38
blob as expert 3 of the winning MoE-V3	gate V3: $+0.00566$ Dice vs its own baseline, 1st of 29 arms (5/5 folds)	same note l.52, l.115; [DOI 10.5281/zenodo.22903668]

Two independent programs — different datasets, architectures, metrics and probability regimes — converge on the same verdict: **null-or-negative alone, first-of-the-field as an initialised expert of a mixture**. That convergence is this paper’s generality argument.

7. Discussion

7.1 Why equal weight per instance does not pay alone, in this regime

The measurements support a coherent reading, offered as a hypothesis grounded in the tables (no gradient-level evidence is collected here). Under **exclusive softmax**, every probability mass given to class c inside instance n is taken from the other 18 classes at the same pixel. The per-instance term weights each instance’s *presence* equally, so its gradient keeps pushing class- c mass onto small instances long after the region-wise losses (CE-ISNS, Dice) have settled — mass that lands as **dispersed coverage**: enough scattered correct pixels to raise thin-class IoU ($+0.98$ traffic light) and contour F1 ($+0.63$), not enough to keep small masks **connected**. The per-class fragment decomposition (§6.4) shows the cost: holes multiply inside ground-truth instances of 16 classes of 19 vs the paired arm B — of 19 of 19 vs the control (person $\times 2.2$), and a holed mask loses instance recall (-4.34 strict), small-instance recall first (T1 $-5.65 > T2 -4.58 > T3 -2.18$) and pedestrian pixel precision (-2.70). The class-level mIoU integrates the scattered gains and the scattered losses into a **null** ($+0.170$, $p = 0.4032$): the term redistributes correctness, it does not create it. On BRATS (sigmoid, independent channels), the same term was also null alone — the competition regime is not the whole story, but the *usage* verdict replicated exactly.

7.2 Where it pays: an initialised expert of a mixture

The mixture arm of the Cityscapes plateau, **MoE-V3-CS**, initialises four experts from the independently trained arms B, D, Dp, G and trains a top-2 gate over 3×3 patches for 80 epochs (control: CE+Dice 80 epochs, B-initialised). Read from the plateau artifacts (master table P3.14, versatility P3.17 — context, §6.1 and §6.5), this arm gives: Δ mIoU **$+0.45$ pt** [**$+0.11$; $+0.82$**], **$p = 0.0066$** vs its control; small-instance recall **T1 $+0.80$** ; traffic light **not degraded** (-0.08 , where the boundary/distmap arms lose 2.3–2.8); maximal damage **-0.53 pt**; fragments 488/image — *below* the control’s 519. The profile is that of an arm putting G’s thin-object speciality to use without carrying its hole-making failure mode; the mechanism — averaging well-initialised experts rather than gate selection — is itself published on BRATS with the same expert: **1st of 29 arms, 5/5 folds** [DOI 10.5281/zenodo.22903668]. **The blob loss is a good expert and a bad generalist** — and initialised mixture-of-experts training is the mechanism that, on both datasets, converts one into the other.

7.3 Practical takeaways

- **Do not ship the blob term alone** in exclusive-softmax semantic segmentation at fixed budget: expect a null mIoU, better thin-object pixel IoU, and degraded pedestrian instance integrity. If pedestrian *detection* feeds a downstream planner, the -4.3 strict-recall points are disqualifying on their own.
- **Do consider it as an expert candidate** for a patch-wise mixture: it is the only arm of the 13 that improves traffic light vs the control, and its speciality (thin objects) is precisely what an expert-initialised mixture puts to use on both datasets (§7.2).
- **If you keep it, pair it with a CC filter**. The fragmentation diagnosis ($+318.8$ components/image) says a post-hoc connected-component consensus veto — the mIoU-neutral $C \oslash B$ of [Cassez 2026b], -18.7% fragments — targets precisely this arm’s failure mode. (Reported as a design implication; the $G \oslash$ -style veto on this arm is not measured in this paper.)

- **The engineering cost is solved:** precomputed instance packs make the term +1–2 %/epoch (against +870 s/epoch naive), parity-tested, with the flip-alignment pitfall documented (§4).

7.4 Relation to Kofler et al.

Nothing here contradicts the published medical result: under the region-based sigmoid recipe the term improved the small-structure metrics it targets [Kofler 2023]. What this paper adds is the **regime boundary**: transposed *exactly* (same algebra, same $\beta = 0.5$, same two-level mean) to exclusive-softmax full-resolution Cityscapes, the term’s promise inverts into the trade-off above. Loss-term transfers across probability regimes should be measured, not assumed — on both of our datasets the *alone* verdict was null while the *expert* verdict was first-of-the-field on BRATS and the only T0 arm of the Cityscapes plateau once inside the mixture.

8. Limitations

1. **The primary endpoint is null.** This is a negative-result paper by design; the contribution is the diagnostic, the port and the usage — not a performance claim.
2. **β not swept.** $\beta = 0.5$ fixed at Kofler’s default ($\alpha:\beta = 2:1$); no λ sweep, same fixed-budget stance as the program’s companion papers. The trade-off could move with β ; that surface is not measured.
3. **Holdout first: 500.** A pre-specified subset of val, identical for all 13 arms and fixed before training — but not the official leaderboard, which was not submitted to.
4. **Holm leaves no arm significant on mIoU in either multiplicity family** (best raw $p = 0.0048 \rightarrow$ Holm 0.058 over the 12-pair master-table family, 0.075 over the program’s 15-pair exploratory family). Program-wide conclusions on mIoU rest on amplitudes and CIs, stated as such.
5. **Budget mismatch with the 13-arm context.** G trains 160 epochs; the context control 80. The paired comparison of G is B (160 epochs); the control pair is always labelled *context*.
6. **Mechanism is a measurement-grounded hypothesis.** The fragmentation decomposition is new and blocking-checked, but no gradient-level or per-layer analysis is performed; §7.1 remains an interpretation.
7. **$n = 3$ seeds.** Sub-0.5-pt effects are underpowered at the seed level; the primary inference is the 500-image paired bootstrap, which probes evaluation-set sampling, not seed variance.
8. **cuDNN/BF16 non-determinism declared.** Two independent forwards of the same checkpoints differ by ≤ 0.01 pt (verdicts identical); bit-exact reruns are not claimed anywhere in the program.
9. **One architecture, two datasets.** ConvNeXt-V2-Base + UPerNet only; Cityscapes + BRATS only. Other heads (Mask2Former-style mask attention re-weights boundaries internally) may react differently.
10. **Instance packs not shipped** (heavy); regenerated bit-identical by the released script, parity verified over all 2 975 packs on 2026-09-26.

9. Conclusion

Given every ground-truth instance the same weight, whatever its pixel count, and full-resolution exclusive-softmax Cityscapes answers: **not worth it alone**. The pre-registered primary endpoint is null ($\Delta\text{mIoU} +0.170$ pt [$-0.233 ; +0.551$], $p = 0.4032$, G vs paired B, 160 epochs, 3 seeds). The term buys exactly what it promises — thin-object pixel coverage (traffic light +0.98, pole +0.58, bicycle +0.57, Holm-significant; Boundary F1 +0.63) — and pays for it in pedestrian instance integrity (strict recall -4.34 , smallest tercile -5.65 , crowd -5.74) and in mask topology (components $\times 1.52$, 16 classes of 19 up vs paired B and 19 of 19 vs the control, person masks $\times 2.2$), which leaves it a **Pareto-dominated specialist** on the program’s 36-endpoint versatility criterion. The same term, as an initialised expert of a patch-wise mixture, is **first of 29 arms on BRATS** and, on the Cityscapes plateau, enters the **only arm that meets the T0 criterion** (MoE-V3-CS, $\Delta\text{mIoU} +0.45$, $p = 0.0066$, maximal damage -0.53 ; tables P3.14/P3.17). Equal weight per instance does not pay alone — it pays as an expert.

Code, configs, per-arm artifacts pointers, tables, figures and regeneration scripts: github.com/guillaume-cassez/cityscapes-blob-loss-kofler. Program companions: Cityscapes distmap [DOI 10.5281/zenodo.21006236], Cityscapes boundary ablation [DOI 10.5281/zenodo.21006393], BRATS MoE-V3 [DOI 10.5281/zenodo.22903668], BRATS boundary [DOI 10.5281/zenodo.22906447], BRATS consensus [DOI 10.5281/zenodo.22904810], BRATS distmap [DOI 10.5281/zenodo.20110976].

Author contributions

Guillaume Cassez (lead author): study design, model training, evaluations and analyses, manuscript writing. **Stanislas Larnier**: formulation of the research questions, methodological guidance, careful reviews of successive drafts of the paper. Both authors approved the final version and the author order.

Appendix A — Runtime and reproducibility

Stage	Hardware	Time	Output
Instance-pack precomputation (one-shot, train split)	8 P-cores, cc3d	22.4 ms/img	.blob.npz per image (blob_glob/counts/csr)
Training 160 ep, 1 seed (arm G)	1 × RTX PRO 6000 96 GB Blackwell	29.2 h (656 s/ep, 31.5 GB)	checkpoints/pilot_fullres_G_ blob_seed{s}/epoch_160.pth
Arm G total (3 seeds)	same	≈ 87 h GPU	final train loss 0.3743 / 0.3735 / 0.3765
Primary bootstrap (harness)	CPU	minutes	results/moe_v3_cs/harness/ table_Gseul_P310.json
Business metrics + fragments (2 arms × 3 seeds)	1 GPU + 6 workers	≈ 2.6 h (12-arm job)	metriques_G_seed*.npz, frag_*_seed*.npy
Tables + figures of this paper	CPU	minutes	papers/paper4/{tables,figures}/

Reproducibility seeds. Training seeds 42, 123, 456 are set globally (PyTorch, NumPy, Python random, CUDA). The bootstrap seed is fixed (20260618, B = 10 000) and shared by every table of the program, so all CIs and p-values are exactly re-derivable from the released per-image artifacts. cuDNN benchmark is left **on** (`deterministic: false`) for training speed; bit-exact GPU reruns are therefore not claimed — the declared reproducibility statement is: same artifacts → same statistics bit-for-bit (verified), and independent forwards of the same checkpoints → agreement within 0.01 pt (measured, 14/14 blocking sanity checks).

Appendix B — Per-class fragmentation, G vs B (components per image)

8-connected components per class per image, 500 images × 3 seeds, paired bootstrap B = 10 000 (seed 20260618), Holm over the 19-class family. Bold = Holm < 0.05. New measurement of this paper (the program had fragments only at arm level before it).

Class	G	B	ctrl	$\Delta(G-B)$	95 % CI	p	Holm(19)
car	117.3	67.9	57.4	+49.4	[+46.5 ; +52.2]	<0.0001	<0.001
pole	176.4	131.1	115.9	+45.4	[+42.9 ; +47.8]	<0.0001	<0.001
vegetation	107.7	63.8	58.8	+43.8	[+41.6 ; +46.1]	<0.0001	<0.001
road	82.4	43.8	47.1	+38.6	[+36.3 ; +40.9]	<0.0001	<0.001
traffic sign	92.3	53.9	43.2	+38.4	[+36.1 ; +40.7]	<0.0001	<0.001
building	116.9	84.4	75.4	+32.5	[+29.6 ; +35.4]	<0.0001	<0.001
sidewalk	106.2	75.4	64.6	+30.8	[+28.7 ; +33.0]	<0.0001	<0.001
person	53.2	24.3	16.4	+28.9	[+27.3 ; +30.5]	<0.0001	<0.001
sky	27.7	23.2	11.9	+4.6	[+3.2 ; +5.9]	<0.0001	<0.001
fence	10.9	8.2	4.9	+2.7	[+2.1 ; +3.4]	<0.0001	<0.001
bicycle	11.8	9.6	6.2	+2.2	[+1.7 ; +2.7]	<0.0001	<0.001
rider	4.1	2.0	1.9	+2.2	[+2.0 ; +2.4]	<0.0001	<0.001
traffic light	6.9	6.0	4.6	+0.9	[+0.6 ; +1.1]	<0.0001	<0.001
truck	0.6	0.3	0.3	+0.2	[+0.1 ; +0.4]	<0.0001	<0.001
motorcycle	0.7	0.6	0.6	+0.0	[-0.0 ; +0.1]	0.4940	1.000
train	0.3	0.2	0.2	+0.0	[-0.0 ; +0.0]	0.3292	1.000
bus	0.5	0.5	0.4	-0.0	[-0.1 ; +0.1]	0.7028	1.000
wall	7.2	7.2	4.0	-0.0	[-0.5 ; +0.5]	0.9496	1.000
terrain	10.9	12.8	5.5	-1.8	[-2.6 ; -1.1]	<0.0001	<0.001

Sanity: the per-class sums equal the arm-level component counts **bit-for-bit** (G 933.9, B 615.1, control 519.3 per image). Source: results/moe_v3_cs/paper4_blob/table_frag_classes_G_vs_B.json.

References

- Cordts *et al.* (2016). *The Cityscapes dataset for semantic urban scene understanding*. CVPR.
- Cassez (2026a). *Distance-Map Auxiliary Regression for Full-Resolution Cityscapes Segmentation*. Zenodo. DOI 10.5281/zenodo.21006236.
- Cassez (2026b). *Boundary Loss Ablation for Full-Resolution Cityscapes Segmentation: When Dice Helps and When It Doesn't*. Zenodo. DOI 10.5281/zenodo.21006393.
- Cassez & Larnier (2026c). *The Gate Does Not Choose: An Expert-Initialised Mixture-of-Experts Outperforms 24 Arms on BRATS 2023*. Zenodo. DOI 10.5281/zenodo.22903668.
- Cassez & Larnier (2026d). *Precision Pays: A Fixed-Weight Boundary Loss Anchors the Precision End (BRATS)*. Zenodo. DOI 10.5281/zenodo.22906447.
- Cassez & Larnier (2026e). *Two Models That Agree Beat the Best of Them Alone: Parameter-Free Connected-Component Consensus (BRATS)*. Zenodo. DOI 10.5281/zenodo.22904810.
- Cassez (2025). *Distance Map Auxiliary Loss for Brain Tumor Segmentation: A Fragment-Centric Study*. Zenodo. DOI 10.5281/zenodo.20110976.
- Isensee *et al.* (2021). *nnU-Net: a self-configuring method for deep learning-based biomedical image segmentation*. Nature Methods 18.
- Kervadec *et al.* (2019). *Boundary loss for highly unbalanced segmentation*. MIDL. arXiv:1812.07032.
- Kofler *et al.* (2023). *Blob loss for biomedical image segmentation*. IPMI. arXiv:2205.08209.
- Lin *et al.* (2017). *Focal loss for dense object detection*. ICCV.
- Milletari *et al.* (2016). *V-Net: fully convolutional neural networks for volumetric medical image segmentation*. 3DV.
- Perazzi *et al.* (2016). *A benchmark dataset and evaluation methodology for video object segmentation*. CVPR.
- Silversmith *et al.* (2019). *cc3d: connected-components labeling for large 3D arrays (software)*. github.com/seunglab/connected-components-3d.
- Woo *et al.* (2023). *ConvNeXt V2: co-designing and scaling ConvNets with masked autoencoders*. CVPR. arXiv:2301.00808.
- Xiao *et al.* (2018). *Unified perceptual parsing for scene understanding*. ECCV. arXiv:1807.10221.

Manuscript — 2026-09-30. Source code, configs, tables, figures and regeneration scripts: github.com/guillaume-cassez/cityscape-blob-loss-kofler. Author: Guillaume Cassez, independent researcher (ORCID 0009-0007-0987-3931), guillaume-cassez.fr — currently looking for ML / computer vision engineering opportunities.